

Can an LLM Judge Replace Human Coders for Measuring Scientific-Evidence Use in Policy? A Feasibility Study and a Cautionary Result

Nitheesha Nakka

Bruce A. Desmarais

Sarah Rajtmajer

Jeffrey J. Harden

Frederick J. Boehmke

June 26, 2026

Abstract

Can we measure, at scale and *without human coders*, whether the scientific evidence that legislators invoke supports the claims they make? We build an end-to-end pipeline that (i) extracts empirical claims from legislators’ press releases, (ii) retrieves candidate peer-reviewed science, and (iii) assigns each (claim, reference) pair a stance (support/refute/silent). The hard step is the third: assessing *support* has historically required scarce expert coders. We test a fully autonomous alternative—a high-capacity large language model (LLM) as an in-context judge providing “gold” labels, independent open-source natural-language-inference (NLI) classifiers as a scalable surrogate, and prediction-powered inference (PPI) / design-based supervised learning (DSL) to combine them. We report this as a *feasibility study*, not a validated measurement, because we deliberately use *no human-coded labels*: the foundations we build on (the alternative-annotator test; DSL/PPI) formally require a trusted human anchor, which we lack, so our estimates are conditional on a single, unvalidated LLM judge. Two results follow, and we read them cautiously. First, our *zero-shot* NLI configuration agrees with the LLM judge only weakly ($\kappa \approx 0.15\text{--}0.21$) and labels most pairs “silent”; anchoring on the LLM judge and applying PPI moves the estimated support rate from a naive 0.21/0.11 (D/R) to 0.71/0.54 “under the LLM judge.” We argue this $\sim 3\times$ gap is too large to attribute cleanly to NLI error: it is equally consistent with a misconfigured NLI baseline and with the LLM judge *over-calling* support—and, because the claims and references were themselves produced by LLMs and our judge is an LLM, with *self-preference* (a documented bias) inflating the result. Self-preference and the absence of a human anchor are first-order threats, not footnotes. Our contribution is therefore a reusable, no-API, no-RA *template* for this kind of measurement, together with a cautionary result about what it can and cannot establish on its own.

1 Introduction

Legislators routinely reach for scientific authority. A bill is “backed by research,” a problem is established with a statistic, an intervention is said to be “proven” to work. In technically complex domains—and few are more complex or fast-moving than artificial intelligence (AI)—these appeals to evidence are both more common and harder to evaluate. Yet a basic question has remained effectively unmeasurable at scale: when a legislator invokes science to support a claim, does the scientific record actually support that claim? Answering it would let us move the study of research

use in policy (Weiss, 1979; Oliver et al., 2022) from case studies and surveys toward systematic, population-level measurement.

The obstacle is not data but *validation*. Legislators’ public communications are abundant and machine-readable, and modern language models can extract the empirical claims they contain and retrieve plausibly relevant scientific literature. The hard step is the last one: deciding whether a retrieved study *supports*, *refutes*, or is *silent* on a given claim. This judgment has traditionally required trained human coders, whose cost and scarcity cap the size and scope of any such study and whose reliability must itself be established. As a practical matter, the “does-the-science-support-the-claim” step is where projects of this kind stall.

We show that this step can be performed *autonomously*, with no human research assistants and no proprietary application programming interfaces (APIs), while preserving statistical validity. Our design rests on three ideas from the recent literature. First, high-capacity large language models (LLMs) now match or exceed trained crowd coders on relevance and stance tasks (Gilardi et al., 2023; Ziems et al., 2024), so an LLM can supply gold-standard labels. Second, because any machine labels are imperfect and naively substituting them for human labels biases downstream estimates (Egami et al., 2023), we combine a small set of high-quality LLM-judge labels with a large set of cheap, independent classifier labels using prediction-powered inference (PPI) (Angelopoulos et al., 2023) and design-based supervised learning (DSL) (Egami et al., 2023). Third, because any single judge can be biased (Zheng et al., 2023), we pair a generative LLM judge with discriminative natural-language-inference (NLI) classifiers that fail in different ways, and we report their agreement explicitly, in the spirit of the alternative-annotator test (Calderon et al., 2025).

We apply the system to a corpus of 7,139 California State Assembly press releases, 303 of which discuss AI, and to a parallel pilot in K–12 education policy, and we report two results cautiously. *First, methodologically*, our *zero-shot* NLI configuration agrees with the LLM judge only weakly ($\kappa \approx 0.15\text{--}0.21$) and labels most pairs “silent,” because strict entailment between an abstract and a press-release claim is usually “neutral” even when the study substantiates the claim’s substance. Anchoring on the LLM judge and applying PPI moves the estimated support rate from a naive 0.21 (Democrats) / 0.11 (Republicans) to 0.71 / 0.54. We resist the tempting reading that “the LLM judge recovers the truth.” A correction this large ($\sim 3\times$) is atypical when a surrogate is competent, and is equally consistent with a misconfigured NLI baseline (zero-shot entailment is a poor proxy for scientific support; serious claim verification uses fine-tuning and rationale selection (Wadden et al., 2020)) and with the LLM judge *over-calling* support. Because the claims and references were produced by LLMs and our judge is an LLM, *self-preference*—the documented tendency of LLM evaluators to favor model-generated text (Panickssery et al., 2024)—could inflate the support rate, and we have no human anchor to rule this out. *Second, substantively and provisionally*, under our judge both parties’ AI claims are supported at broadly similar (and fairly high) rates, so the large partisan gap implied by the naive NLI estimate does not survive; what is robust is that Democrats discuss AI far more often and more intensively than Republicans, not any difference in the support for their claims.

Our contributions are deliberately modest. (1) We provide a reusable, end-to-end, fully autonomous *template*—claim extraction, evidence retrieval, and support assessment—that runs with no human coders and no paid APIs, released as open-source software, that others can apply and, crucially, validate against a human standard. (2) We show that a popular shortcut—scoring scientific support with off-the-shelf zero-shot NLI—can be badly miscalibrated for press-release claims, a cautionary result for the growing practice of measuring political text with such tools. (3) We are explicit about what an autonomous, single-model, self-judging pipeline *cannot* establish on its own (validity against truth), and we specify what a valid version would require. We do *not* claim a validated estimate of evidence use; the substantive patterns we report are provisional and condi-

tional on an unvalidated judge. The remainder of the paper describes the related literatures (§2), the data (§3), our methods (§4), and our results (§5), before discussing implications and limitations (§6–§7).

2 Related work

Evidence use in policymaking. A long line of scholarship asks how, and how much, research evidence informs policy. Weiss (1979) distinguished instrumental, conceptual, and symbolic uses of research; more recent work emphasizes the organizational and relational conditions under which evidence is taken up (Oliver et al., 2022; Bogenschneider et al., 2019). This literature is overwhelmingly qualitative or survey-based and rarely measures, at scale, whether the evidence a policymaker invokes actually supports the claim being made. We provide such a measure, turning a normative concern about “evidence use” into a quantity that can be estimated from text.

Text as data and legislative communication. Press releases are a workhorse corpus for studying how legislators present themselves and frame policy (Grimmer, 2013), and computational text analysis is now standard in political science (Grimmer and Stewart, 2013; Monroe et al., 2008). We build on this tradition but shift the unit of analysis from words or topics to *claims* and their relationship to an external scientific record—a comparatively underdeveloped target that requires linking political text to the research literature.

LLMs as annotators. A rapidly growing body of work shows that large language models can match or exceed trained crowd workers on text-annotation tasks including relevance and stance detection, at a fraction of the cost and with higher consistency (Gilardi et al., 2023; Alizadeh et al., 2023; Ziems et al., 2024). This makes an autonomous coding pipeline plausible. But naive adoption is dangerous: model labels contain non-random error, and substituting them for human labels biases downstream estimates and breaks inference even at high accuracy (Egami et al., 2023). Two complementary responses have emerged. First, statistical corrections—design-based supervised learning (Egami et al., 2023) and prediction-powered inference (Angelopoulos et al., 2023)—recover valid estimates by combining many surrogate labels with a small gold set. Second, validation protocols such as the alternative-annotator test (Calderon et al., 2025) specify *when* a model may be treated as a replacement annotator. We use both, and we add a practical ingredient: the gold labels themselves are produced by a high-capacity LLM judge rather than by humans, with independent NLI models as the surrogate, so the entire procedure is autonomous.

LLM-as-judge and its pitfalls. Using one model to evaluate another is now common (Zheng et al., 2023), but LLM judges exhibit position, verbosity, and self-preference biases and can inherit the validity problems of whatever they are calibrated against. The standard mitigation is to avoid relying on a single judge. Our design follows this prescription by pairing a generative LLM judge with discriminative NLI classifiers that fail in different ways, and by reporting cross-method agreement explicitly rather than assuming it.

Scientific claim verification. Determining whether a document supports a claim is the subject of scientific fact-checking and claim-verification research, typically framed as natural-language inference over claim–evidence pairs (Williams et al., 2018; Yin et al., 2019; Wadden et al., 2020). Koneru et al. (2024) study whether LLMs can discern evidence for scientific hypotheses in the social sciences specifically. Our stance step adopts this framing but embeds it in a measurement pipeline

whose output is a debiased *population* quantity (the party-level support rate), not a per-pair prediction, which is what shifts the validity burden onto PPI/DSL.

Contribution. Relative to this work we contribute (i) an end-to-end, fully autonomous pipeline—claim extraction, evidence retrieval, and support assessment—that needs neither human RAs nor proprietary APIs; (ii) a concrete demonstration that strict-entailment NLI systematically under-detects scientific support, so that an LLM judge is not a luxury but is required for an unbiased estimate; and (iii) substantive findings on evidence use in AI policymaking that this machinery makes estimable for the first time.

3 Data

Corpus. Our primary corpus is the universe of press releases issued by members of the California State Assembly and posted to their official sites, harvested in 2025. After de-duplication it comprises **7,139** documents—5,249 from Democratic and 1,890 from Republican members—spanning roughly 2010 to 2025. California is a deliberate choice: it is the most active and visible state in AI policymaking (National Conference of State Legislatures, 2024), and its lopsided partisan balance (roughly three-quarters Democratic) is itself substantively relevant and is handled explicitly in estimation.

AI-related subset. We flag documents containing AI-related language (e.g. “artificial intelligence,” “machine learning,” “algorithmic,” “facial recognition,” “deepfake,” “large language model”) and extract the AI-bearing sentences. **303** documents (254 Democratic, 49 Republican) contain at least one such statement, for a total of **910** AI statements. The partisan asymmetry is large on every margin: 4.8% of Democratic documents discuss AI versus 2.6% of Republican documents, and Democratic AI documents average 3.3 AI statements each versus 1.6 for Republicans. Consistent with prior descriptions of this corpus, a Fightin’-words analysis (Monroe et al., 2008) shows Democrats emphasizing generative AI, health, and data, and Republicans emphasizing accountability, fairness, and enforcement.

Claims and retrieved evidence. From the AI-related documents the pipeline extracts empirical claims and retrieves candidate scientific references (§4.1). Our analysis set is the resulting collection of **621** (claim, reference) pairs for which a reference with an abstract was retrieved (292 from Democratic, 329 from Republican claims). The retrieved literature is recent and broadly open: reference years concentrate after 2015 and rise sharply through 2024, the single most common venue is arXiv, and 87% of retrieved references are open-access (83% for Democratic claims, 90% for Republican). These descriptive facts establish that a substantial, accessible scientific literature exists for the claims legislators make about AI; whether that literature *supports* the claims is the question we take up in §5.

Education-policy pilot. To show that the pipeline is not specific to AI, we replicate its front end on a separate corpus of California Assembly press releases on K–12 education: 5,131 documents (4,279 Democratic, 852 Republican), of which 611 (11.9%) are education-related. Partisan attention is again highly asymmetric: Democrats issued 80 releases on social-emotional learning and student mental health versus 1 for Republicans, and 38 on school safety versus 4. The education pilot is descriptive here and serves to demonstrate domain transfer.

Validation sample. Our autonomous validation requires gold labels for a sample of claim–reference pairs. These are produced by the LLM judge (§4.2) on a party-stratified sample drawn from the 621-pair analysis set; no human-coded labels are used. We report results as a function of the gold-sample size, since it governs the width of the prediction-powered confidence intervals.

4 Methods

Our goal is a measurement procedure that maps each (claim, scientific reference) pair to a stance label—SUPPORT, REFUTE, MIXED, or SILENT (irrelevant or not addressed)—and that aggregates these labels into valid, party-level estimates of evidence use, *without* requiring human coders. We first describe the upstream pipeline that produces the claim–reference pairs (§4.1), then the autonomous validation framework that is our central contribution (§4.2).

4.1 The policy-to-science pipeline

The pipeline converts a policymaking document into a set of (*claim* $\rightarrow$ *candidate evidence*) records in three stages.

1. **Claim extraction.** Each press release is processed by a large language model prompted to extract up to ten empirical claims “for which, if you were writing a scientific paper, you would need a reference.” Output is structured (press-release id, claim number, claim text), isolating the verifiable propositions from rhetorical or procedural text (Grimmer, 2013).
2. **Evidence retrieval.** For each claim, a web-connected LLM returns up to ten candidate scientific references as structured JSON (title, abstract, authors, year, venue, DOI, URL, open-access flag), with de-duplication by DOI (else title+year). This stage answers “what science is *relevant* to this claim,” a task at which current models are strong but which does not, by itself, establish *support*.
3. **Support assessment.** Each (claim, reference) pair is assigned a stance label. This is the step that previously required expert coders and is the focus of §4.2.

The substantive object of interest is the *support rate*: among relevant claim–reference pairs, the share for which the scientific record supports the legislator’s claim, and how this varies by party.

4.2 Autonomous validation: LLM judge + NLI surrogate + PPI

Validators. We treat stance assignment as an annotation problem and replace human annotators with two kinds of automated *validators* that fail in different ways, so that their agreement is informative (Zheng et al., 2023).

- **LLM judge.** A high-capacity instruction-tuned model (here Claude Opus 4.8, queried in context with no fine-tuning and no API key) reads the claim and the reference title+abstract and returns a stance label with a brief rationale. LLMs of this class match or exceed trained crowd coders on relevance and stance tasks (Gilardi et al., 2023; Ziemis et al., 2024).
- **NLI classifiers.** Two open-source natural-language-inference models with different architectures and training data—DeBERTa-v3-base-mnli-fever-anli and bart-large-mnli—score each pair with the reference as premise and the claim as hypothesis. We map ENTAILMENT $\rightarrow$ SUPPORT, CONTRADICTION $\rightarrow$ REFUTE, and NEUTRAL $\rightarrow$ SILENT (Yin et al., 2019;

Laurer et al., 2024). NLI is the standard discriminative approach to scientific claim verification (Wadden et al., 2020; Koneru et al., 2024).

- **Dependency-free baselines.** TF-IDF cosine relevance and a negation-aware lexical rule provide validators that run with no models at all, for continuous-integration and no-dependency reproduction.

Agreement as autonomous reliability. For any pair of validators we report percent agreement and Cohen’s κ over the shared, non-abstaining items. This plays the role that inter-annotator reliability plays in human coding. Following the *alternative-annotator test* (Calderon et al., 2025), a validator is admissible as a stand-in only if it is at least as good a predictor of the reference standard as the reference annotators are of one another; agreement that is high *among methodologically similar* validators but low *across* methods is diagnostic of a shared blind spot rather than of correctness.

Valid estimation under imperfect labels. Plugging machine labels directly into a downstream analysis biases estimates and invalidates confidence intervals, even at 80–90% label accuracy (Egami et al., 2023). We therefore treat the cheap, scalable NLI labels as a *surrogate* available for all N pairs, and the LLM-judge labels as *gold* available for a sample of $n \ll N$ pairs, and combine them with prediction-powered inference (Angelopoulos et al., 2023). For the binary support indicator, the PPI point estimate of the population support rate is

$$\hat{\theta}_{\text{PPI}} = \underbrace{\frac{1}{N} \sum_{i=1}^N f(x_i)}_{\text{surrogate mean (all pairs)}} + \underbrace{\frac{1}{n} \sum_{j=1}^n (y_j - f(x_j))}_{\text{rectifier (gold sample)}}, \quad (1)$$

where $f(x_i) \in \{0, 1\}$ is the NLI “support” prediction and $y_j \in \{0, 1\}$ is the LLM-judge gold label, with variance $\widehat{\text{Var}} = \text{Var}(f(x))/N + \text{Var}(y - f(x))/n$ and a Wald 95% interval $\hat{\theta}_{\text{PPI}} \pm 1.96 \widehat{\text{Var}}^{1/2}$. When the rectifier is non-zero—i.e. when the surrogate is biased relative to the judge—PPI corrects the naive surrogate mean while retaining the efficiency gains of labeling all N pairs. Equation (1) is the prediction-powered special case of the design-based supervised learning estimator (Egami et al., 2023); both are formally valid when the gold labels are sampled by a known design, and both reduce to the human-coded estimate as $n \rightarrow N$.

Why this is “autonomous”—and what that costs. Every component runs locally with no human annotation and no paid API: the LLM judge stands in for the human RA, the NLI surrogate supplies scale, and PPI combines them. But “autonomous” is not free. The PPI/DSL guarantee (Angelopoulos et al., 2023; Egami et al., 2023) holds only when the anchor labels are *trusted ground truth* sampled by a known design, and the alternative-annotator test (Calderon et al., 2025) that licenses replacing human coders is *defined* as a comparison to multiple human annotators. We have neither: our anchor is a single, unvalidated LLM. Our confidence intervals are therefore *conditional* on that judge—they quantify sampling of the judge’s labels, not error relative to truth—and should be read as illustrative, not as validated population inference. We treat this as the central limitation of the design, not a detail, and return to it in §6. The moment a small set of *human-verified* labels becomes available, it can replace (or augment) the LLM-judge anchor to obtain a conventionally valid estimate.




Figure 1: Democrats engage with AI more often (a) and more intensively (b) than Republicans in California Assembly press releases.

Threats specific to an all-LLM loop. Two further issues follow from doing everything with language models. First, *self-preference*: LLM judges favor model-generated and low-perplexity text and can recognize their own outputs (Panickssery et al., 2024; Zheng et al., 2023); because the claims and the retrieved references in our pipeline are themselves LLM-produced, the judge may rate them as supportive partly for that reason, biasing the support rate upward. Second, *retrieval bias*: the retrieval step seeks *relevant* references and performs no deliberately disconfirming search, so a pair scored “support” is supported *by what an LLM chose to surface*, and the refute rate is a lower bound. We de-duplicate identical (claim, reference) pairs before analysis, and we flag that some retrieved “references” are bill text or government pages rather than peer-reviewed science, which a faithful stance step should treat as out of scope.

Implementation. The pipeline and validator are released at <https://github.com/bdesmarais/ai-policy-science-use> (`scripts/core/autonomous_validation.py`); all results below are reproducible from the committed reference outputs with a single command and no API key.

5 Results

5.1 Engagement with AI is large and partisan

Democrats engage with AI far more than Republicans on every margin (Figure 1). Democratic members mention AI in 4.8% of their press releases versus 2.6% for Republicans, and when they do, they say more about it: Democratic AI documents contain 3.3 AI statements on average versus 1.6 for Republicans. Across the 303 AI documents the pipeline retrieves a substantial and accessible body of candidate evidence—614 claim–reference pairs with abstracts, 87% of them open-access, drawn from venues led by arXiv and concentrated in the years since 2015. The raw material for evidence-based AI policy clearly exists; the question is whether legislators’ claims are borne out by it.

Table 1: Cross-validator agreement. The two discriminative NLI models agree with each other but not with the generative LLM judge.

Validator pair	n	% agreement	Cohen’s κ
Claude judge $\sim$ DeBERTa-NLI	120	43.3	0.21
Claude judge $\sim$ BART-NLI	120	39.2	0.15
DeBERTa-NLI $\sim$ BART-NLI	614	86.5	0.66




Figure 2: The generative judge and the discriminative NLI classifiers agree poorly with each other ($\kappa \approx 0.15$ – 0.21), while the two NLI classifiers agree well ($\kappa = 0.66$). This flags a shared method property—bias or shared error—rather than establishing which validator is correct.

5.2 Generative and discriminative validators disagree by method

A first result is that the two *kinds* of validator disagree sharply (Table 1, Figure 2). The two NLI classifiers, which share the discriminative entailment paradigm, agree strongly with each other ($\kappa = 0.66$ over all 614 pairs), but each agrees only weakly with the generative LLM judge ($\kappa = 0.21$ and 0.15 over the 120 gold pairs). The mechanism is that strict textual entailment between a study’s abstract and a press-release claim is almost always classified “neutral”—hence SILENT—even when the study substantiates the claim’s substance, because the abstract rarely restates the claim in entailing language: our zero-shot NLI setup labels 76% of pairs SILENT, the LLM judge only 23%. Two cautions are essential. (i) This is *our zero-shot NLI configuration*, not a fair representative of automated claim verification; strong systems on tasks like SciFact (Wadden et al., 2020) fine-tune and select rationale sentences and do not equate “non-entailment” with “no support,” so we cannot claim that NLI *as a class* under-detects support. (ii) High *within-method* agreement reflects a shared property of the two NLI models—which could be shared conservatism *or* shared error—and does not establish which validator is correct; absent a human anchor, neither does the disagreement. We report the disagreement as a flag, not as proof that the LLM judge is right (Calderon et al., 2025).

Table 2: Share of claims “supported,” by party. The PPI estimate is anchored on a single, unvalidated LLM “gold” ($n = 120$) and is therefore *conditional on that judge*, not validated against humans.

Party	Naive NLI	PPI [95% CI] (cond.)	LLM-judge raw
Democratic	0.21	0.71 [0.57, 0.84]	0.70
Republican	0.11	0.54 [0.41, 0.67]	0.57

5.3 Under the LLM judge, support rates look high—but this is a conditional estimate

Relative to the naive NLI surrogate (0.21 for Democratic, 0.11 for Republican claims), anchoring on the LLM judge’s 120 gold labels and applying PPI (Eq. 1) roughly triples the estimates, to **0.71** (95% CI [0.57, 0.84]) for Democrats and **0.54** ([0.41, 0.67]) for Republicans (Table 2, Figure 3); the judge’s raw rates (0.70 and 0.57) are similar, as they must be, since the judge *is* the anchor. We deliberately do not headline this as “most claims are supported.” A roughly threefold gap between surrogate and judge is far larger than is typical when a surrogate is competent (in the PPI/DSL literature surrogates are usually 80–90% accurate and corrections are modest), and we cannot, without a human anchor, attribute it to NLI error rather than to the judge *over*-calling support—a live possibility given self-preference in an all-LLM loop (§4). The level is also inflated by construction: retrieval is relevance-seeking with no disconfirming search, and some “supporting” references are bill text or government pages rather than peer-reviewed studies. We therefore read the level as “high under our judge,” not as an established fact about evidence use.

What is more robust is *comparative* and *negative*. The large partisan gap implied by the naive NLI estimate (Democratic claims nearly twice as likely to be “supported”) does not survive a better judge: under the LLM judge the rates are 0.71 vs. 0.54 with overlapping intervals. The apparent partisan asymmetry in evidentiary quality is thus an artifact of a poorly-configured instrument; the asymmetry that is robust across instruments is the *volume and intensity* of AI engagement (Fig. 1), not the support for the claims.

5.4 Domain transfer: education policy

The pipeline’s front end transfers to a different polarized domain. On a separate corpus of 5,131 education-policy press releases, 611 (11.9%) are education-related, with the same kind of partisan asymmetry in attention (Figure 4): Democrats issued 80 releases on social-emotional learning and student mental health versus 1 for Republicans, and 38 on school safety versus 4. These asymmetries in *which* topics receive attention motivate extending the full support-assessment analysis to education in future work.

6 Discussion

A competent judge is necessary, but not sufficient. The most actionable lesson is negative: our off-the-shelf zero-shot NLI configuration is not a safe stand-alone measure of whether science supports a policy claim. It labels genuine support “neutral” because abstracts seldom restate a press-release claim in entailing language, and the downward bias is large enough to flip a substantive conclusion. But the corollary is *not* “therefore trust the generative judge.” A competent judge appears necessary, yet it is not sufficient: to license replacing human coders it must itself be




Figure 3: Estimated share of claims “supported.” The naive NLI surrogate (gray) differs sharply from the single-LLM judge anchor; PPI (colored, 95% CIs, *conditional on that unvalidated judge*) roughly triples the estimate and narrows the partisan gap. The large gap is consistent with a mis-configured NLI baseline and/or the judge over-calling support; it is not validated against humans.

validated against a human standard (Calderon et al., 2025), which we have not done. We present the NLI failure as a warning about a popular shortcut, not as proof that the generative judge is right. (Strong claim verification, e.g. on SciFact (Wadden et al., 2020), fine-tunes and selects rationale sentences; our zero-shot setup is not a fair representative of that line of work.)

The part we did *not* solve: validity without humans. PPI (Angelopoulos et al., 2023) and DSL (Egami et al., 2023) let researchers combine abundant imperfect labels with a small set of *trusted* labels. The tempting extension—that the trusted labels can come from an LLM rather than a human—is exactly what we cannot justify: both methods, and the alternative-annotator test, are defined relative to human ground truth, so substituting an unvalidated model for the human anchor removes the guarantee rather than preserving it. Our intervals are therefore conditional on the judge. We demonstrate the *mechanics* of an autonomous surrogate-plus-judge-plus-PPI pipeline; we do not demonstrate that its output is valid.

The all-LLM loop and self-preference (a first-order threat). The deepest problem is circularity. Our claims and references are LLM-produced and our judge is an LLM, and LLM evaluators are documented to recognize and favor model-generated, low-perplexity text (Panickssery et al., 2024); more broadly, fluent model output should not be mistaken for knowledge (Bender et al., 2021). A judge that rates LLM-surfaced evidence as “supportive” may be exhibiting self-preference rather than assessing science, and the $\sim 3\times$ naive-to-PPI jump is the magnitude such a bias could produce. Without a non-self anchor we cannot separate “the judge corrects the surrogate” from “the judge agrees with other models.” This threatens the central estimate directly; it is not a footnote.

Diagnose across methods, not within—but a flag is not a validation. One genuine lesson survives: within-method agreement can be an illusion of correctness (our two NLI models agreed at $\kappa = 0.66$ while sharing a bias), so cross-method agreement is a more informative diagnostic. But cross-method agreement only *flags* disagreement; it cannot tell us which side is right. Only




Figure 4: Partisan asymmetry in education-policy attention (counts of press releases), illustrating that the pipeline’s front end generalizes beyond AI.

comparison to a trusted standard can do that. We recommend reporting cross-method agreement routinely while stressing that it does not substitute for a human anchor.

What this does and does not say about AI policymaking. We are deliberately cautious. The one robust, instrument-independent pattern is that Democrats engage with AI far more, and more intensively, than Republicans. A useful *negative* result is that the apparent partisan gap in evidence *quality* from the naive measure does not survive a better judge. We do *not* claim that “most claims are supported”: that level (≈ 0.5 – 0.7 under our judge) is conditional on an unvalidated, possibly self-preferring judge, is inflated by relevance-only retrieval and by bill-text/government “references,” and could move substantially under human coding. The contribution is a template plus a cautionary tale, not a measurement of how well legislators’ AI claims are backed by science.

A template, honestly scoped. The design—cheap open-source surrogate, a small high-quality gold set, PPI/DSL, no API, no RA—is a useful scaffold that lowers cost and access barriers, but its value is realized only when the gold is trustworthy; with an LLM gold it is a feasibility scaffold awaiting validation. Reproducibility (one command, no key) is necessary but not sufficient: perfectly reproducing an unvalidated measurement reproduces its bias.

7 Limitations and Conclusion

Limitations. The binding limitation is that *we have no human validation*. Our “gold” is a single, unvalidated LLM judge, so the alternative-annotator test (Calderon et al., 2025) is not actually run and the PPI/DSL guarantees (Angelopoulos et al., 2023; Egami et al., 2023) do not formally apply; every support-rate number is conditional on that judge. Compounding this, the judge and the upstream claim/reference steps are all LLMs, so *self-preference* (Panickssery et al., 2024) may inflate the support rate and we cannot rule it out without a non-self anchor. Beyond these: (3) *stance is a proxy*—“support” read from a title and a ≤ 600 -character abstract is a coarse stand-in for a methods-aware reading; (4) *retrieval bounds the ceiling*—we score only relevance-retrieved evidence with no disconfirming search, so the support rate is upward-biased and the refute rate a lower bound, and some “references” are bill text or government pages, not science; (5) *the gold set*

is small ($n = 120$) and from one model, so the intervals are wide *and* understate uncertainty about the gold itself; (6) *strategic text*—press releases are advocacy, so “support” for a chosen claim is not balanced evidence use; (7) *scope*—one chamber of one heavily-Democratic state, in English.

What a valid version would require. The path from this feasibility scaffold to a valid measurement is concrete: (i) a *human*-coded anchor of roughly 100–200 claim–reference pairs, sampled by a known design, against which to run the alt-test and to seed DSL/PPI; (ii) a properly-configured claim-verification baseline (fine-tuned, rationale-based, in the SciFact line (Wadden et al., 2020)) rather than zero-shot entailment, so that “NLI” is given a fair test; (iii) a *non-self* judge—a different model family, or human coders—to break the all-LLM circularity and quantify self-preference; (iv) full-text, methods-aware stance rather than abstract-only; and (v) disconfirming retrieval. None of these is exotic; together they would convert the conditional estimates here into validated ones, and the released scaffold is built so that adding a human anchor is a drop-in change.

Conclusion. We asked whether claim–evidence *support* can be measured at scale with no human coders. Our answer is a qualified “not yet, but here is a usable scaffold.” We can build the autonomous pipeline—claim extraction, retrieval, and an LLM judge anchoring a cheap NLI surrogate via PPI—and run it end-to-end on open models with no API key; doing so yields a genuine cautionary result (off-the-shelf zero-shot NLI is badly miscalibrated for this task) and a robust descriptive finding (asymmetric AI *engagement* across parties). But the headline support rates are conditional on an unvalidated, possibly self-preferring judge, and we decline to present them as established. The contribution is a reproducible template and an honest map of what stands between it and a valid measurement, released so that others can supply the missing human anchor and put it to the test.

Reproducibility and ethics. All code, the claim–reference data, the LLM-judge gold labels, and the figure scripts are available at <https://github.com/bdesmarais/ai-policy-science-use>; every number in this paper is reproducible with a single command and no paid API. The corpus consists of public communications by elected officials acting in their official capacity; we make no claims about private individuals. Because automated “evidence scoring” of political speech could be misused to lend false objectivity to partisan judgments, we stress that our estimates are measurements with stated uncertainty and known biases, not verdicts, and that the cross-method diagnostics and human-anchor option are integral to using the tool responsibly.

AI-assistance disclosure

This paper is, by design, an AI-heavy artifact, and we disclose that explicitly because it bears on how the results should be weighed. The stance “gold” labels were produced by Claude Opus 4.8 acting as an in-context judge (no fine-tuning, no API key); the scalable surrogate labels were produced by open-source NLI models (DeBERTa- and BART-MNLI); and the same model assisted in drafting this manuscript. Every cited reference was checked against a primary source (see `paper/reviews/refcheck.md` in the repository); none was found to be fabricated. Two critical self-reviews—one methodological, one on the paper as an AI-produced artifact—shaped this revision and are included in the repository. We emphasize that computational reproducibility does *not* establish validity: re-running an unvalidated, single-model measurement reproduces its biases exactly. All design decisions and claims are the authors’. Earlier pipeline development was supported by NSF awards 2318460 and 2148215.

References

- Alizadeh, M. et al. (2023). Open-source large language models outperform crowd workers and approach ChatGPT in text-annotation tasks. *arXiv preprint arXiv:2307.02179*.
- Angelopoulos, A. N., Bates, S., Fannjiang, C., Jordan, M. I., and Zrnic, T. (2023). Prediction-powered inference. *Science*, 382(6671):669–674.
- Bender, E. M., Gebru, T., McMillan-Major, A., and Shmitchell, S. (2021). On the dangers of stochastic parrots: Can language models be too big? In *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency (FAccT)*, pages 610–623.
- Bogensneider, K., Day, E., and Parrott, E. (2019). Revisiting theory on research use: Turning to policymakers for fresh insights. *American Psychologist*, 74(7):778.
- Calderon, N., Reichart, R., and Dror, R. (2025). The alternative annotator test for LLM-as-a-judge: How to statistically justify replacing human annotators with LLMs. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics*. arXiv:2501.10970.
- Egami, N., Hinck, M., Stewart, B. M., and Wei, H. (2023). Using imperfect surrogates for downstream inference: Design-based supervised learning for social science applications of large language models. In *Advances in Neural Information Processing Systems (NeurIPS)*. arXiv:2306.04746.
- Gilardi, F., Alizadeh, M., and Kubli, M. (2023). Chatgpt outperforms crowd workers for text-annotation tasks. *Proceedings of the National Academy of Sciences*, 120(30):e2305016120.
- Grimmer, J. (2013). *Representational Style in Congress: What Legislators Say and Why It Matters*. Cambridge University Press.
- Grimmer, J. and Stewart, B. M. (2013). Text as data: The promise and pitfalls of automatic content analysis methods for political texts. *Political Analysis*, 21(3):267–297.
- Koneru, S., Wu, J., and Rajtmajer, S. (2024). Can large language models discern evidence for scientific hypotheses? case studies in the social sciences. In *Proceedings of LREC-COLING 2024*, pages 2787–2797.
- Laurer, M., van Atteveldt, W., Casas, A., and Welbers, K. (2024). Less annotating, more classifying: Addressing the data scarcity issue of supervised machine learning with deep transfer learning and BERT-NLI. *Political Analysis*, 32(1):84–100.
- Monroe, B. L., Colaresi, M. P., and Quinn, K. M. (2008). Fightin’ words: Lexical feature selection and evaluation for identifying the content of political conflict. *Political Analysis*, 16(4):372–403.
- National Conference of State Legislatures (2024). Artificial intelligence 2024 legislation. <https://www.ncsl.org>.
- Oliver, K., Hopkins, A., Boaz, A., Guillot-Wright, S., and Cairney, P. (2022). What works to promote research-policy engagement? *Evidence & Policy*, 18(4):691–713.
- Panickssery, A., Bowman, S. R., and Feng, S. (2024). LLM evaluators recognize and favor their own generations. In *Advances in Neural Information Processing Systems (NeurIPS)*. arXiv:2404.13076.

- Wadden, D., Lin, S., Lo, K., Wang, L. L., van Zuylen, M., Cohan, A., and Hajishirzi, H. (2020). Fact or fiction: Verifying scientific claims. In *Proceedings of EMNLP*.
- Weiss, C. H. (1979). The many meanings of research utilization. *Public Administration Review*, 39(5):426–431.
- Williams, A., Nangia, N., and Bowman, S. R. (2018). A broad-coverage challenge corpus for sentence understanding through inference. In *Proceedings of NAACL-HLT*.
- Yin, W., Hay, J., and Roth, D. (2019). Benchmarking zero-shot text classification: Datasets, evaluation and entailment approach. In *Proceedings of EMNLP-IJCNLP*.
- Zheng, L., Chiang, W.-L., Sheng, Y., et al. (2023). Judging LLM-as-a-judge with MT-Bench and chatbot arena. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- Ziems, C., Held, W., Shaikh, O., Chen, J., Zhang, Z., and Yang, D. (2024). Can large language models transform computational social science? *Computational Linguistics*, 50(1):237–291.